

SEUPD@CLEF: Team retrix on Source Retrieval for Scientific Web Claims

Fabio Baldan¹, Giancarlo Padoan¹, Alvisè Garberino¹, Marco Tessari¹ and Davide Donati¹

¹University of Padua, Italy

Abstract

We present the results of the retrix group for the CheckThat! 2026 - Task 1, in which the goal is to identify the scientific paper implicitly referenced in a social media post. The input consists of short, noisy and usually informal posts that mention scientific findings without providing explicit links or explicit citations.

Our approach is designed to operate in a multilingual setting, as the dataset includes posts in English, German, and French. To handle this, we used multi-lingual language models.

DA FINIRE CON PIPELINE FINALEEEEEEE DA FINIRE CON PIPELINE FINALEEEEEEE DA FINIRE CON PIPELINE FINALEEEEEEE DA FINIRE CON PIPELINE FINALEEEEEEE

Keywords

Fact Checking, Information Retrieval, Scientific Document Retrieval, Social media fact verification, CLEF, Neural Re-ranking, Cross Encoder

1. Introduction

The rapid spread of scientific claims through social media creates a difficult fact-checking scenario: posts often report, simplify, or sensationalize research findings while omitting the original source. In many cases, the cited study is not linked explicitly with a URL or a bibliographic reference, but only mentioned indirectly through fragments such as the reported finding, the venue, the institution, or the authors. As a consequence, verifying such claims requires first identifying the scientific publication that is being implicitly referenced.

This problem is the focus of Task 1 of the CLEF 2026 CheckThat! Lab, *Source Retrieval for Scientific Web Claims*. Given a social media post containing a scientific claim and an implicit reference to a paper, the goal is to retrieve the correct publication from a fixed collection of 10,000 candidate papers; the task can be naturally framed as an information retrieval problem in which short, noisy posts must be matched against much longer and structurally richer scientific documents composed of titles, abstracts, venues, and author information.

The task is challenging for several reasons. First, the linguistic gap between social media and scientific writing is substantial: posts are short, informal, and often contain paraphrases, slogans, hashtags, or incomplete statements, whereas scientific papers use more precise and technical language. Second, the reference to the source paper is implicit rather than explicit, so the retrieval system must infer the connection from partial semantic evidence rather than from direct citations. Third, relevant clues may appear in different forms, such as mentions of a scientific result, a medical outcome, a research institution, or a publication venue, requiring the system to combine lexical and semantic evidence effectively.

An additional challenge is multilinguality; the official query sets include English, German, and French posts, while the candidate pool remains shared across languages. Therefore, the retrieval system must handle both cross-style and cross-language variation: it has to connect informal user-generated content

“Search Engines”, course at the master degree in “Computer Engineering”, Department of Information Engineering, and at the master degree in “Data Science”, Department of Mathematics “Tullio Levi-Civita”, University of Padua, Italy. Academic Year 2025/2026

✉ fabio.baldan.1@studenti.unipd.it (F. Baldan); giancarlo.padoan@studenti.unipd.it (G. Padoan); alvisè.garberino@studenti.unipd.it (A. Garberino); marco.tessari.8@studenti.unipd.it (M. Tessari); davide.donati@studenti.unipd.it (D. Donati)



© 2026 Copyright for this paper by its authors.

Use permitted under Creative Commons License Attribution 4.0 International (CC BY 4.0).

with formal scientific text, and it must do so consistently across multiple languages. This makes the task particularly suitable for investigating retrieval architectures that can balance lexical matching, semantic similarity, and robust ranking strategies.

In this work, we study retrieval approaches for this task and analyze how different components contribute to the identification of the correct source paper; rather than focusing only on final ranking quality, we also consider the broader retrieval setting in which recall, robustness across languages, and compatibility with downstream re-ranking all play a central role.

The paper is organized as follows: Section 2 discusses the most relevant research directions related to this task; Section 3 describes our approach; Section 4 explains our experimental setup; Section 5 discusses our main findings; finally, Section 6 draws some conclusions and outlooks for future work.

2. Related Work

The task addressed in this paper lies at the intersection of fact-checking and information retrieval. A natural starting point is lexical retrieval, where probabilistic ranking models such as BM25 remain strong baselines thanks to their efficiency, interpretability, and robustness in candidate generation [?]. In the setting of scientific claim source retrieval, lexical matching is especially useful when a post mentions recognizable entities or terminology that also appears in the title, abstract, venue, or author list of the referenced paper.

Neural retrieval methods have been proposed to mitigate the lexical mismatch between queries and documents by learning dense semantic representations; more broadly, neural information retrieval has shifted attention from exact token overlap to representation learning and learned matching functions [?]. In particular, dense dual-encoder architectures such as Dense Passage Retrieval (DPR) demonstrated that learned query and document embeddings can be highly effective for first-stage retrieval [?]. This line of work is especially relevant to our task because social media posts often paraphrase scientific findings instead of reusing the exact wording found in the original publication.

Between initial retrieval and final ranking, many modern systems adopt multi-stage pipelines that trade off efficiency and effectiveness. Late-interaction models such as ColBERT preserve fine-grained token-level matching while remaining substantially more scalable than fully pairwise neural scoring over the whole collection [?]. Likewise, transformer-based rerankers have become a common second-stage component for refining an initial candidate set with stronger query-document interactions [?]. These approaches motivate retrieval architectures that first generate a manageable set of candidates and then apply more expensive relevance estimation only where it is most useful.

Our work is also related to the broader literature on automatic claim verification and, more specifically, to the CheckThat! shared tasks, which have addressed the detection, verification, and analysis of claims in social media across several editions [? ? ?]; however, source retrieval for scientific web claims differs from standard verification tasks in an important way: the system is not asked to classify a claim as true or false, but to recover the original scientific paper implicitly referenced in the post and for this reason, the task is closely connected to evidence retrieval, while also requiring robust matching between informal multilingual posts and formal scientific documents.

3. Methodology

Describe the methodology you have adopted, the architecture of your system, your workflow, etc.

We started by implementing a basic pipeline using BM25 as baseline, using Apache Lucene as the underlying framework. We then experimented with different stop lists, filters and weightings to improve the MRR@5 score, used as final score in the CheckThat! Lab. We then proceeded to experiment with multi-stage pipelines based on neural models, which is the final structure of our project.

We implemented a three-stage pipeline for our information retrieval system, consisting of:

- **Query expansion:** We expanded the original query with additional context to bridge the gap between informal and scientific language.




Figure 1: Three-stage pipeline for information retrieval.

- **Retrieval:** We retrieved relevant documents from the corpus using the expanded query.
- **Reranking:** We reranked the retrieved documents based on their relevance to the original query.

For each of these stages, we experimented with different models, choosing the ones that performed better in relation to the others.

3.1. Query Expansion

3.2. Retrieval

3.3. Reranking

3.4. Evaluation

4. Experimental Setup

Describe the experimental setup, i.e.

- used collections
- evaluation measures
- url to git repository and its organization
- hardware used for experiments
- ...

5. Results and Discussion

Provide a summary of the performance on the previous year dataset.

Discuss the results and any relevant issues.

6. Conclusions and Future Work

Provide a summary of what are the main achievements and findings.

Discuss future work, e.g. what you may try next and/or how your approach could be further developed.

7. Group Members Contribution

Jane Doe Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Table 1
Frequency of Special Characters

Non-English or Math	Frequency	Comments
Ø	1 in 1,000	For Swedish names
π	1 in 5	Common in math
\$	4 in 5	Used in business
Ψ_1^2	1 in 40,000	Unexplained usage

John Doe Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Jane Doe Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

John Doe Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Jane Doe Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

John Doe Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

8. Misc [TO BE REMOVED]

8.1. Tex Files

Put your $\text{\LaTeX}$ files into the `section` folder as shown in the examples above.

8.2. Figures

Put your figures into the `figure` folder and put the caption under the figure. Example of reference to Figure 2.

8.3. Tables

Put the caption above the table. Example of reference to Table 1.

See the `booktab` packaged documentation for further options.



Figure 2: 1907 Franklin Model D roadster. Photograph by Harris & Ewing, Inc. [Public domain], via Wikimedia Commons. (<https://goo.gl/VLCRBB>).

8.4. Bibliography

Example of citations:

- name: ?]
- parenthesis: [?]

An initial list of references is provided in the files `bibliography.bib` and `proceedings.bib` that you can expand.

See the `natbib` packaged documentation for further options.

8.5. Acronyms

Use the

`\ac{acronym}`

command to insert acronyms, eg. *Average Precision* (*AP*). The command will expand the acronym the first time it is used.

An initial list of acronyms is provided in the file `acronyms.tex` that you can expand.

See the `acronym` packaged documentation for further options.